
CS165 Final Report

Kevin Zhang

December 20, 2016

1 FEATURE OVERVIEW

Almost all of the basic features requested of this project have been implemented. The last features remaining at the moment are nested-loop joins and updates. The list of features includes:

- Creates - create a database, table, column, or index.
- Inserts - insert a row into any table.
- Selects - select data using a range operator on any column.
- Fetches - fetch data after selecting.
- Loads - load data from files into the database.
- Indexes - create indexes over columns for faster selects.


In addition, deletion in B+ trees has been implemented and tested to work (but not incorporated into updates yet).

Code has been organized according to project components; the Makefile used looks in several pre-defined folders for source and header files. Header files mirror the layout of the source files. In addition, a `data` folder stores all database metadata and values (persisted upon server shutdown and loaded upon server startup). This folder (along with subfolders) is automatically created when needed. A `catalog` file at the root contains database metadata.

Extensive logging has been implemented; when turned on, all queries, server responses, client messages, and more are printed to stdout. In addition, after execution of each query and on startup, the current database and client context are printed to help debug.

2 SYSTEM OVERVIEW

The diagram below details how queries are handled by the server.



Green nodes indicate sections the user interacts with, red nodes indicate data sources, and purple nodes indicate server functions. On a client request, the following actions are taken:

1. Request is forwarded from server handler to query parser.
2. Query parser constructs a `DbOperator` object and passes it to the executor.
3. Executor interacts with the database to build a `Result` object if necessary.
4. Executor returns a string response to the server.
5. Server responds to the client if appropriate.


The `DbOperator` object contains all metadata required to execute queries; a `ClientContext` object is used to store all `Result` objects created from queries.

3 CREATES

Handling create queries is fairly straightforward; on request, corresponding objects are created. For a database, this is a `Db` object; for a table, a `Table` object; for a column, a `Column` object; and for an index, an `Index` object. The `Db`, `Table`, `Column` objects contain straightforward fields describing the contents of each, along with pointers to payloads and contained objects.

`Index` objects are a bit more complicated; these contain a `IndexType` identifying the `Index` as a B+ tree or column and a `clustered` descriptor. `Index` objects additionally contain pointers to their respective `Column` objects (for ease of data lookup) and a `IndexObject`, which is either a `BTreeUNode` (unclustered B+ tree), `BTreeCNode` (clustered B+ tree), `ColumnIndex` (containing integer arrays of values and indices), or `NULL`, signifying a clustered, sorted index (no additional data is needed to maintain this type of index).

`Index` objects are stored in the corresponding `Table`. Upon creation of a clustered index, no data is stored (it is assumed that clustered indices will be declared before any data is inserted into the database). However, upon creation of unclustered indexes, all data is inserted into the indexes. Thus, queries can immediately make use of the indexes (after insertion completes, of course).



4 INSERTS


Inserts without indexes are easy; we append to the end of the corresponding `Column` objects in the given `Table`. Handling `Index` objects is harder.

The system completes the following process whenever inserting data:

1. Search for clustered indexes on the table.
2. If one is found, insert into the clustered index first (no more than one clustered index may exist). Retrieve the indices of the inserted values, then shift all values in each column one by one (to avoid unnecessary random access) as necessary to insert new data.

3. If none are found, append all values to the end of the table, then insert the new (value, index) pairs into each of the unclustered indices.

The B+ tree implementation is critical for this portion of the project (and probably the most time-consuming). The optimal parameters for B+ trees have yet to be determined; **an experiment will be conducted soon to determine optimal fanout/node size.**



5 SELECTS

Selects without indexes are also easy; we simply scan through the entire column and append to a results array contained in the corresponding `ClientContext` object. When we have a usable `Index`, we search through the `Index` to find the smallest value greater than or equal to the minimum of a select query. Depending on the type of index, we then:

- Clustered B+ Tree: search for the maximum, then create an array containing all values between the two indexes found.
- Clustered Sorted: binary search for the maximum, then create an array containing all values between the two indexes found.
- Unclustered B+ Tree: iterate forward through the leaves until we find a value greater than the maximum, and accumulate an array of encountered indices.
- Unclustered Sorted: iterate forward through the array until we find a value greater than the maximum, and accumulate an array of encountered indices.

For batched selects, this behavior is slightly modified:

- Clustered indexes: iterate across select queries first, building result arrays one by one.
- Unclustered indexes: iterate down the index and build result arrays along the way.

For clustered indexes, we avoid leapfrogging between result arrays and instead write all values to each result array first. This avoids some degree of random access. For unclustered indexes, we avoid scanning the same data more than once and instead accumulate values as we go, avoiding multiple data scans but incurring costs from leapfrogging across result arrays instead.